

This project focuses on an image search problem: finding a specific face in an image or photograph. This is a problem that arises in many video surveillance applications (e.g., detecting a well-known terrorist's face in a crowd at an airline terminal) as well as in searching large collections of digital images (e.g., finding all photos you took five years ago that contain your grandmother). In this assignment, we will focus on finding George P. Burdell, whose mugshot is shown in Figure 1, in an image that contains a crowd of faces, such as the sample scene in Figure 2. The faces in the scene, including George's, may be upright or may be rotated by 90, 180, or 270 degrees.




Figure 1: Mugshot of George P. Burdell

Disclaimer: Any resemblance to actual persons, roommates, or relatives is unintentional and purely coincidental.



Figure 2: Sample Crowd of Faces

We would like this task to be performed in real-time, so the computational and storage requirements must be kept to a minimum. We are also concerned with functional correctness of the algorithm (i.e., getting the correct answer), since we do not want to generate frequent false alarms while monitoring crowded public spaces, or worse, miss an important match!

Description: George will appear exactly once (possibly rotated by 90, 180 or 270 degrees) in a random spot in a crowded scene. None of the faces in the scene will overlap with each other and all faces fit completely within the boundaries of the crowd (no partial faces hanging off the edges).

The faces that are not George will differ from him structurally (e.g., wear glasses) and/or by having different color features (e.g., a blue hat or green skin). Your task is to find the one face that exactly matches George's mugshot, but rotated either 0, 90, 180, or 270 degrees.

The crowded scene is a 64x64 image array of cells; each cell is one of eleven colors whose codes are given in the color palette shown in Figure 3. Each face fits in an 12x12 region of cells. The features of the face will always have color codes in the range 1, 2, ..., 8. The background (non-face parts) of the image will always use colors whose codes are 9, 10, or 11.




Figure 3: Color Palette

The image array is provided as input to the program as a linearized array of the cells in row-column order. The first element of the array (location 0) represents the color of the first cell in the first row. This is followed by the second cell (location 1) in that row, etc. The last cell of the first row (location 63) is followed by the first cell of the second row (location 64). This way of linearizing a two dimensional array is called *row-column mapping*. The color code (1-11) is packed in an unsigned byte integer for each cell.

Strategy: Unlike many “function only” programming tasks where a solution can be quickly envisioned and implemented, this task requires a different strategy.

1. Before writing any code, reflect on the task requirements and constraints. *Mentally* explore different approaches and algorithms, considering their potential performance and costs. The metrics of merit are **static code length**, **dynamic execution time**, and **storage requirements**. There are often trade-offs between these parameters. Sometimes *back of the envelope* calculations (e.g., how many comparisons will be performed) can help illuminate the potential of an approach.
2. Once a promising approach is at hand, an HLL implementation can deepen its understanding. Adding loop counters, print statements, or better still, automatic execution analysis tools (e.g., gprof) can give a more accurate estimate of performance and efficiency. HLL implementation (in C for this project) is more flexible and convenient for exploring the solution space. The final version will be completed in machine assembly where changes, big and small, are much more costly. This version is submitted for P1-1.
3. Once a final C version is selected, it time to “be the compiler” and see how it translates to MIPS assembly. This is an opportunity to see how HLL constructs are supported on a machine platform (at the ISA level). This level requires the greatest programming effort; but it also uncovers many new opportunities to increase performance and efficiency. While assembly coding is less common in system design, it will let you see how the “rubber hits the road”.

P1-1 High Level Language Implementation:

In this section, the first two steps described above will be completed. It's fine to start with a simple implementation that produces an answer; this will help deepen your understanding. Then experiment with your best mentally-formed ideas for a better performing solution. Use any instrumentation techniques that help to assess how much work is being done. Each hour spent exploring here will cut many hours from the assembly level coding exercise.

You should use the simple shell C program that is provided `P1-1-shell.c` to allow you to read in a linearized image array. The shell program includes a reader function `Load_Mem()` that loads

the values from a text file. Rename the shell file to P1-1.c and modify it by adding your image search code. Since generating example crowd images is complex, it is best to fire up Misasim, generate a crowd, step forward until the crowd is written in memory, and then dump memory to a file. Your C program should print the array index of the middle red pixel on the top of George's hat and the middle pixel on George's shirt. For example, for the crowd shown in Figure 4 (below), it should print this location as hat: 2731, shirt: 2742, the index of the pixels 42 rows down and 43 and 54 pixels across, respectively. (This test case is provided on the project website as crowd-2731-2742.txt.)

Be sure that your completed assignment can read in an arbitrary crowd, find George, and correctly print his location (hat and shirt positions) since this is how you will receive points for this part of the project. Note: you will not be graded for your C implementation's performance. Only its accuracy and "good programming style" will be considered. Your MIPS implementation might not even use the same algorithm; although it's much easier for you if it does.

When have completed the assignment, submit the single file P1-1.c to the class T-square website. You do not need to submit data files. Although it is good practice to employ a header file (e.g., P1-1.h) for declarations, external variables, etc., in this project please include this information at the beginning of your submitted program file. In order for your solution to be properly received and graded, there are a few requirements.

1. The file must be named P1-1.c.
2. Your submitted file should compile and execute on an arbitrary crowd (produced from Misasim). It should contain the following print statement (unchanged)

```
printf("George is located at: hat pixel %d, shirt pixel %d.\n", HatLoc, ShirtLoc);
```

to print out two integers which locate George. The command line parameters should not be altered from those used in the shell program. Your program must compile and run with gcc under Linux.
3. Your solution must be properly uploaded to the class T-square site before the scheduled due date, **9:00pm on Wednesday, 2 October 2013.**

P1-2 Assembly Level Implementation: In this part of the project, you will write the performance-focused assembly program that solves the Find Tumbling George puzzle. A shell program (P1-2-shell.asm) is provided to get you started. *Your solution must not change the crowd image array (do not write over the memory containing the crowd).*

In this version, execution performance and cost is often equally important. The assessment of your submission will include functional accuracy during 100 trials and performance and efficiency. The code size, dynamic execution length, and operand storage requirements are scored empirically, relative to a baseline solution. *The dynamic instruction length metric is the maximum of the baseline metric and the average dynamic instruction length of the five fastest student submissions.*

Your score will be determined through the following equation:

$$PercentCredit = 2 - \frac{Metric_{Your Program}}{Metric_{Baseline Program}}$$

Percent Credit is then used to determine the number of points for the corresponding points category. Important note: while the total score for each part can exceed 100%, especially bad performance can earn *negative credit*. The sum of the combined performance metrics scores (code size, execution length, and storage) will not be less than zero points. **Finally, the performance scores will be reduced by 10% for each incorrect trial (out of 100 trials). You cannot earn performance credit if your implementation fails ten or more of the 100 trials.**

Library Routines: There are three library routines (accessible via the `swi` instruction).

SWI 570: Create Tumbling Crowd: This routine initializes memory beginning at the specified base address (e.g., `Array`). It sets each byte of the 4096 byte array to the corresponding cell's color code in the 64x64 crowd image array. The color codes are defined in the palette in Figure 3. INPUTS: \$1 should contain the base address of the 1024 word (4096 byte) array already allocated in memory. OUTPUTS: none.

SWI 552: Highlight Position: This routine allows you to specify an offset into the Crowd array and it draws a white outline around the cell at that offset. Cells that have been highlighted previously in the trace are drawn with a gray outline to allow you to visualize which positions your code has visited. INPUTS: \$2 should contain an offset into the Crowd array. OUTPUTS: none. *This is intended to help you debug your code; be sure to remove calls to this software interrupt before you hand in your final submission, since it will contribute to your instruction count.*

SWI 571: Locate Tumbling George: This routine allows you to specify the position of the middle red pixel at the top of George's hat and the middle pixel of his shirt to indicate the location where George has been found at a certain orientation.

INPUTS: \$2 should contain two packed numbers: in the upper 16 bits, the hat pixel position and in the lower 16 bits, the shirt pixel position. Each pixel position must be a number between 0 and 4095, inclusive. This answer is used by an automatic grader to check the correctness of your code.

OUTPUTS: \$3 gives the correct answer. You can use this to validate your answer during testing. If you call `swi 571` more than once in your code, only the first answer that you provide will be recorded. The visualization will draw small yellow boxes at the pixel locations you provide and a larger 12x12 yellow box centered at these locations that you provide. It will also draw small green boxes at the pixel positions of the correct answer and a larger 12x12 green box showing George's actual location. If your answer is correct, the green boxes should completely cover your yellow boxes. For example, in Figure 4, the correct answer is 0xAAB0AB6 (2731 decimal in top 16 bits; 2742 decimal in lower 16 bits) and when `swi 571` is called with this value in \$2, the green boxes appear as shown in Figure 4.



Figure 4: Crowd with Correct Answer: 2731 hat, 2742 shirt

In order for your solution to be properly received and graded, there are a few requirements.

1. The file must be named `P1-2.asm`.
2. Your program must produce and store the correct values in memory in the statically allocated memory.
3. Your program must call SWI 571 to report an answer and return to the operating system via the `jr` instruction. *Programs that include infinite loops or produce simulator warnings or errors will receive zero credit.*
4. Your solution must be properly uploaded to the submission site before the scheduled due date, **9:00pm on Wednesday, 16 October 2009**. **There will be a one hour grace period, after which no submitted material will be accepted.**

Implementation Evaluation:

In this project, the functional implementation of Find Tumbling George in C (**P1-1**) will be evaluated on whether the correct answer is computed. Although proper coding technique (e.g., using the proper data types, operations, control mechanisms, etc.) will be evaluated, parametric performance will not be considered.

The performance implementation of the Find George program in MIPS assembly (**P1-2**) will be evaluated both in terms of correctness and performance. The correctness evaluation employs the same criterion described for the functional implementation. The performance criteria include static code size (# of instructions), dynamic instruction length (# executed instructions), and storage requirements (in this example, number of registers used). All of these metrics are used to determine the quality of the overall implementation.

Once a candidate algorithm is selected, an implementation is created, debugged, tested, and tuned. In MIPS assembly language, small changes to the implementation can have a large impact

on overall execution performance. Often tradeoffs arise between static code size, dynamic execution length, and operand storage requirements. Creative approaches and a thorough understanding of the algorithm and programming mechanisms will often yield impressive results.

One final note on design strategy: while optimizing MIPS assembly language might, at times, seem like a job best left to automated processes (i.e., compilers), it underscores a key element of engineering. Almost all engineering problems require multidimensional evaluation of alternative design approaches. Knowledge and creativity are critical to effective problem solving.

Project Grading: The project grade will be determined as follows:

<i>part</i>	<i>description</i>	<i>due date 9:00pm</i>	<i>percent</i>
P1-1	Find George (C code)	Wed., 2 October 2013	25
P1-2	Find George (MIPS assembly)	Wed., 16 October 2013	
	correct operation, technique & style		25
	static code size		15
	dynamic execution length		25
	operand storage requirements		10
	<i>Total</i>		100

Good luck!